

Write your first Kubernetes  
controller in less than 3 hours

---


# Introduction

---

# Federico Paolinelli



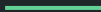
- Telco Network Team @ Red Hat
- All things networking and kubernetes
- MetalLB maintainer



# Francesco Romani



- Telco Compute Team @ Red Hat
- Kubernetes tuning/enhancement for low latency
- Kubernetes SIG-node maintainer



## Workshop structure

- Topic introduction
- Practice
- Practice review

[github.com/goworkshops/k8s-example-kubedredger](https://github.com/goworkshops/k8s-example-kubedredger)

# github.com/goworkshops/k8s-example-kubedredger

github.com/goworkshops/k8s-example-kubedredger/tree/exercise-1

goworkshops / k8s-example-kubedredger

Code Issues Pull requests Actions Projects Wiki Security Insights Settings

k8s-example-kubedredger Public

Edit Pins Watch 0 Fork 0 Star 0

exercise-1 7 Branches 0 Tags Go to file

This branch is 2 commits ahead of main.

Contribute

fe...	Add the first exercise file	6417961 · 7 hours ago	3 Commits
.devcontainer	Initial import	8 hours ago	
.github/workflows	Initial import	8 hours ago	
.dockerignore	Initial import	8 hours ago	
.gitignore	Initial import	8 hours ago	
.golangci.yml	Initial import	8 hours ago	
EXERCISE.md	Add the first exercise file	7 hours ago	

About

Sample kubernetes controller for Go workshop

Activity

Custom properties

0 stars

0 watching

0 forks

Report repository

Releases

No releases published

[Create a new release](#)

Packages


# The Kubernetes model

---

# Kubernetes deals with objects

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: frontend
spec:
  replicas: 3
  selector:
    matchLabels:
      tier: frontend
  template:
    metadata:
      labels:
        tier: frontend
    spec:
      containers:
        - name: php-redis
          image: gcr.io/google_samples/gb-frontend:v3
```

```
kubectl apply -f replica.yaml
```






# Kubernetes deals with objects

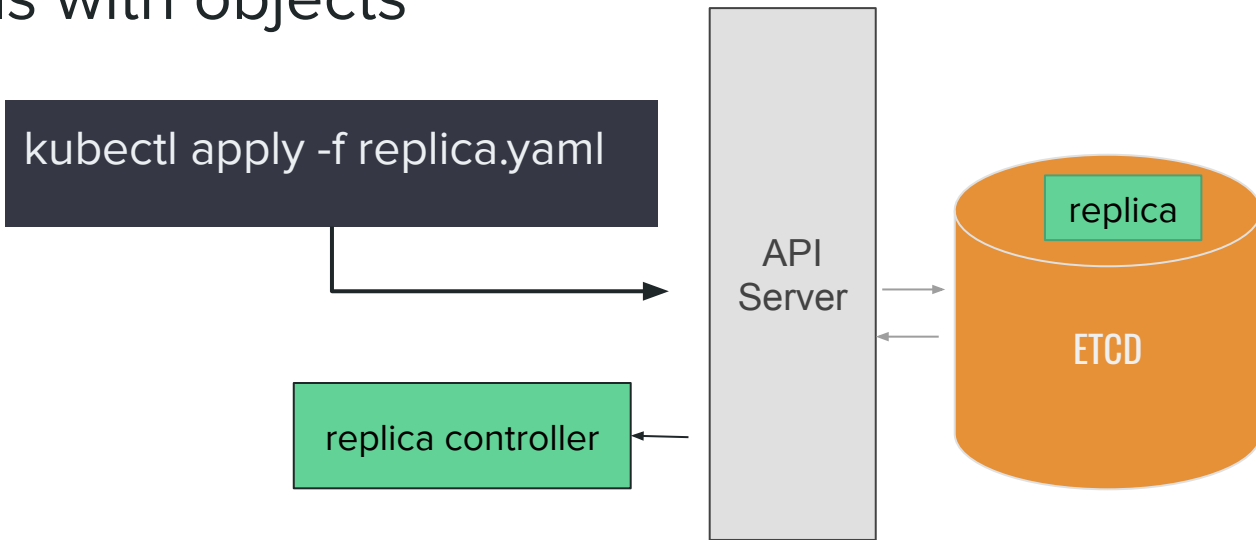
```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: frontend
spec:
  replicas: 3
  selector:
    matchLabels:
      tier: frontend
  template:
    metadata:
      labels:
        tier: frontend
    spec:
      containers:
        - name: php-redis
          image: gcr.io/google_samples/gb-frontend:v3
```

kubectl apply -f replica.yaml




# Kubernetes deals with objects

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: frontend
spec:
  replicas: 3
  selector:
    matchLabels:
      tier: frontend
  template:
    metadata:
      labels:
        tier: frontend
    spec:
      containers:
        - name: php-redis
          image: gcr.io/google_samples/gb-frontend:v3
```



# Declarative over imperative

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: frontend
spec:
  replicas: 3
  selector:
    matchLabels:
      tier: frontend
  template:
    metadata:
      labels:
        tier: frontend
    spec:
      containers:
        - name: php-redis
          image: gcr.io/google_samples/gb-frontend:v3
```




Pod 1

Pod 2

Pod 3



# Declarative over imperative

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: frontend
spec:
  replicas: 2
  selector:
    matchLabels:
      tier: frontend
  template:
    metadata:
      labels:
        tier: frontend
    spec:
      containers:
        - name: php-redis
          image: gcr.io/google_samples/gb-frontend:v3
```



A Kubernetes cluster is a dynamic system

---



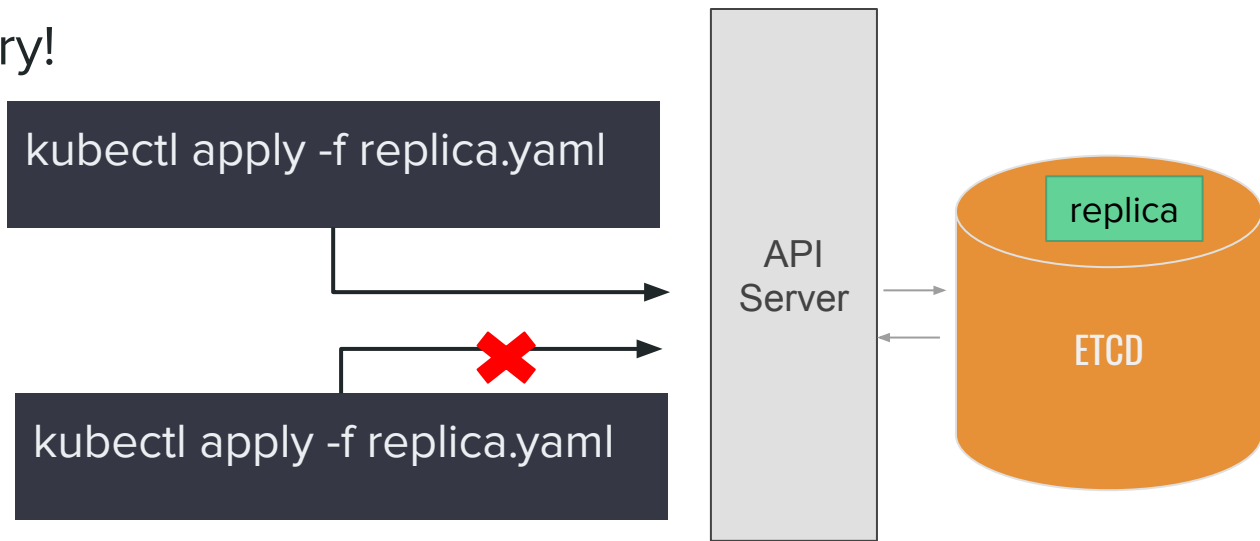


# Complications

<https://flic.kr/p/HCWbEi>

# Knock knock! Race condition! Who's there?

- Kubernetes is read intensive
- Concurrency is handled optimistically
- Conflicts are resolved with failures
- The client must retry!





# Missed Events!

Your controller can be restarted at any time

The network is not reliable

The whole node can crash!




Edge driven - Level driven

---


# Edge Driven vs Level Driven

number of pods




Let's consider the variation of the number of pods in a deployment.

# Edge Driven



When driven by edges, the behaviour depends on the variation of the data observed at edge location.

# Leven Driven



With levels, we care about the level, in our case the state of the system at a given time.

## Edge Driven vs Level Driven?

- Handling the variation might seem easier
- **Leven driven is more robust**
- The reconciliation logic is more complex

# The theory

---

# Objects

---



# Kubernetes objects taxonomy

`/apis/batch/v1/namespaces/F00/jobs`

# Kubernetes objects taxonomy

- Group



/apis/**batch**/v1/namespaces/F00/jobs

# Kubernetes objects taxonomy

- Group
- Version



/apis/batch/**v1**/namespaces/F00/jobs

# Kubernetes objects taxonomy

- Group
- Version
- Resource



`/apis/batch/v1/namespaces/F00/jobs`

All the objects share the same structure

---

## Common traits

```
type Pod struct {  
    metav1.TypeMeta  
    metav1.ObjectMeta  
    Spec    PodSpec  
    Status  PodStatus  
}
```

## TypeMeta

```
type TypeMeta struct {  
    Kind      string  
    APIVersion string  
}
```

## ObjectMeta


```
type ObjectMeta struct {  
    Name           string  
    GenerateName   string  
    Namespace      string  
    SelfLink       string  
    UID            types.UID  
    ResourceVersion string  
    Generation     int64  
    .  
    .
```

# Spec and Status


```
type Pod struct {  
    metav1.TypeMeta  
    metav1.ObjectMeta  
    Spec    PodSpec  
    Status  PodStatus  
}
```




Spec is the desired state, status the actual state




Spec is the desired state, status the actual state



Spec is the desired state, status the actual state



Spec is the desired state, status the actual state



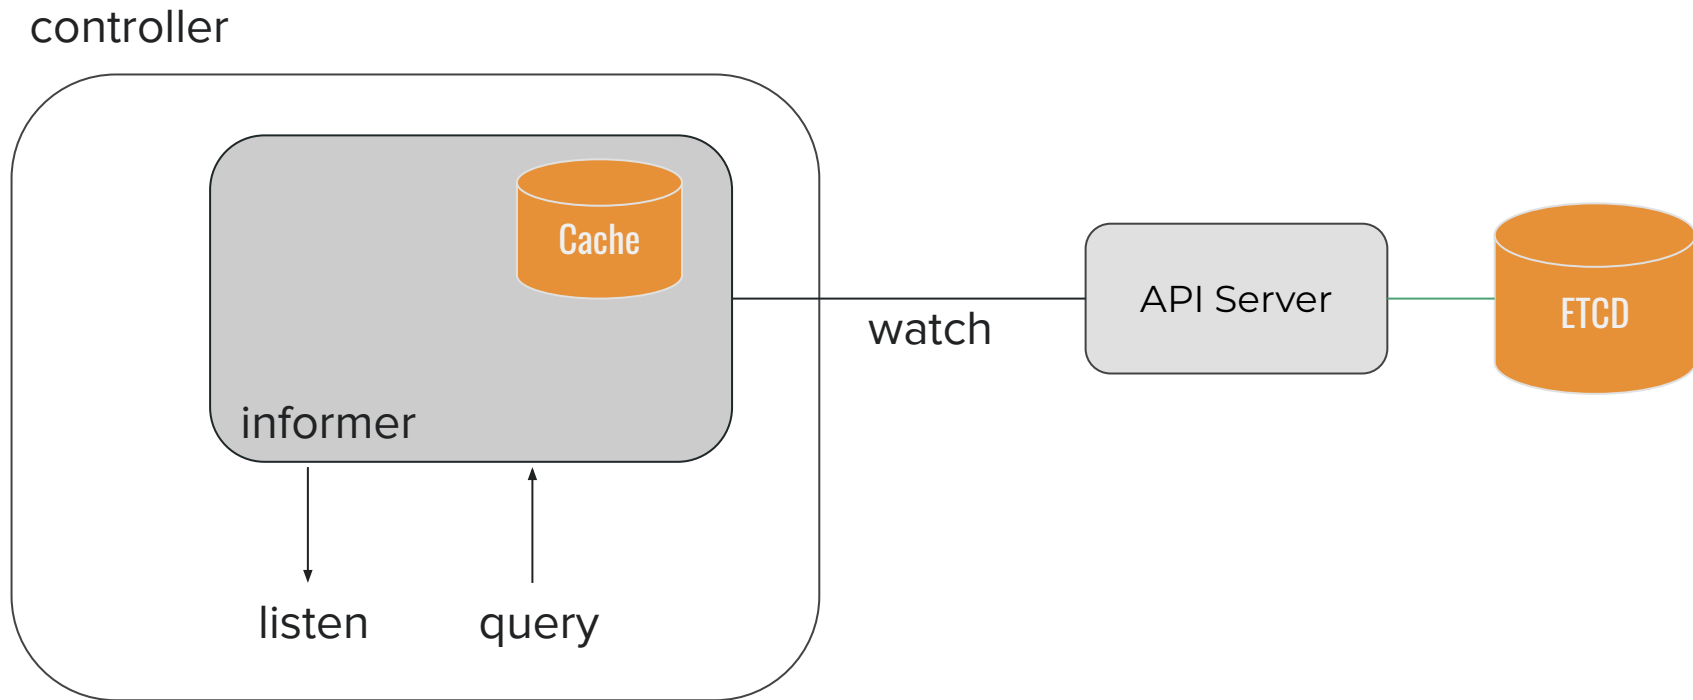
# The building blocks

---

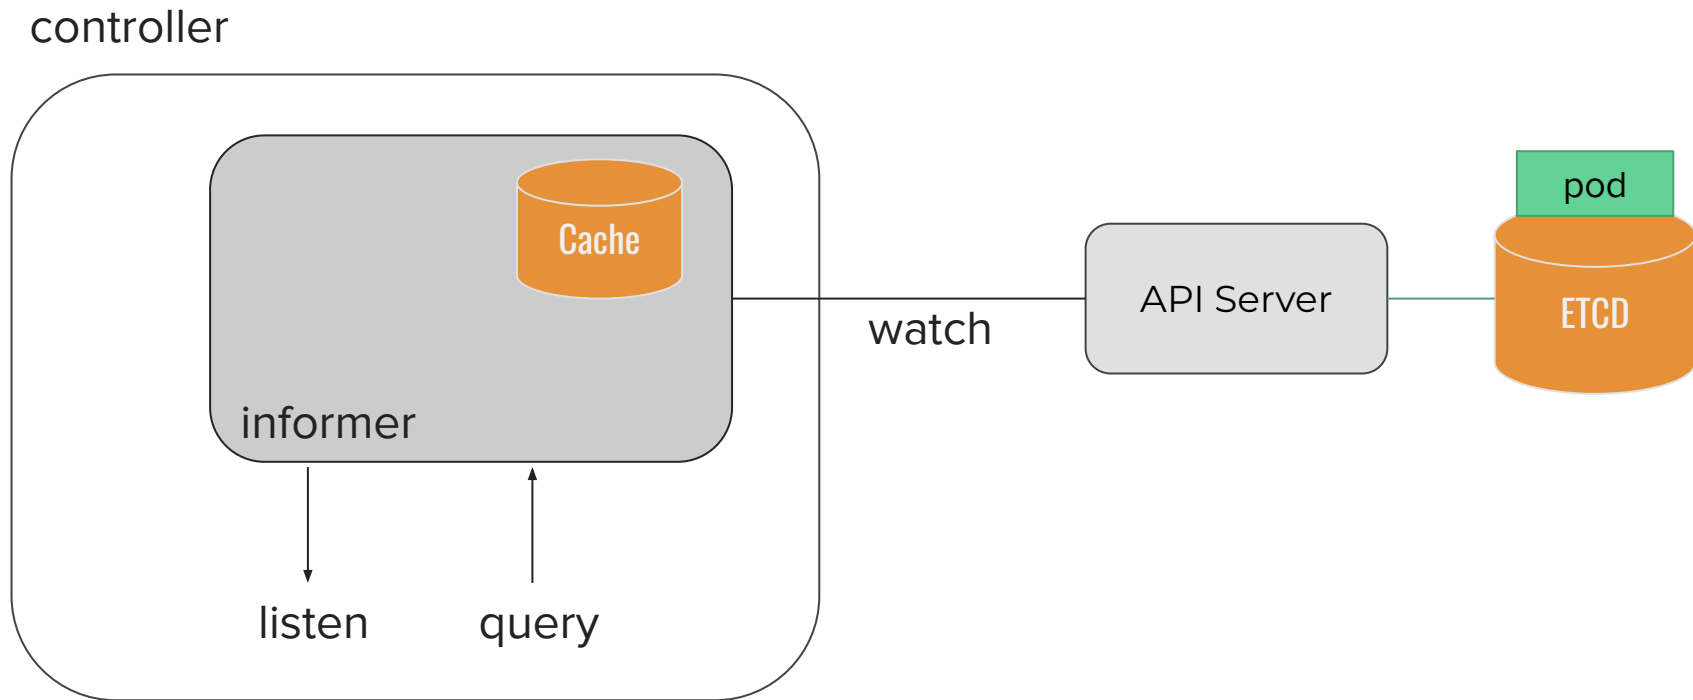
# Informers

- Contain a local cache
- Can watch for resources
- Handles reconnections
- Can notify a client

# Informers

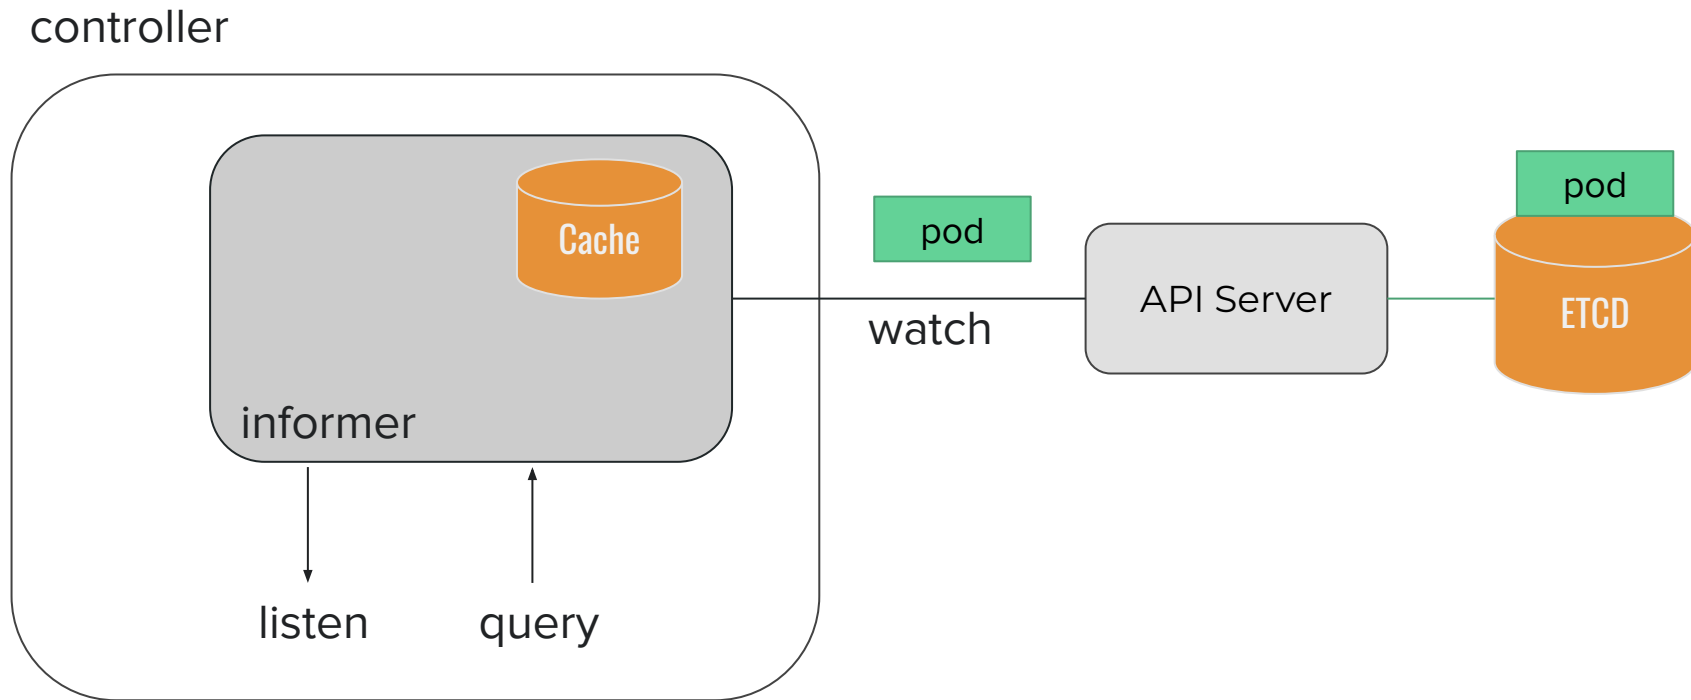


# Informers

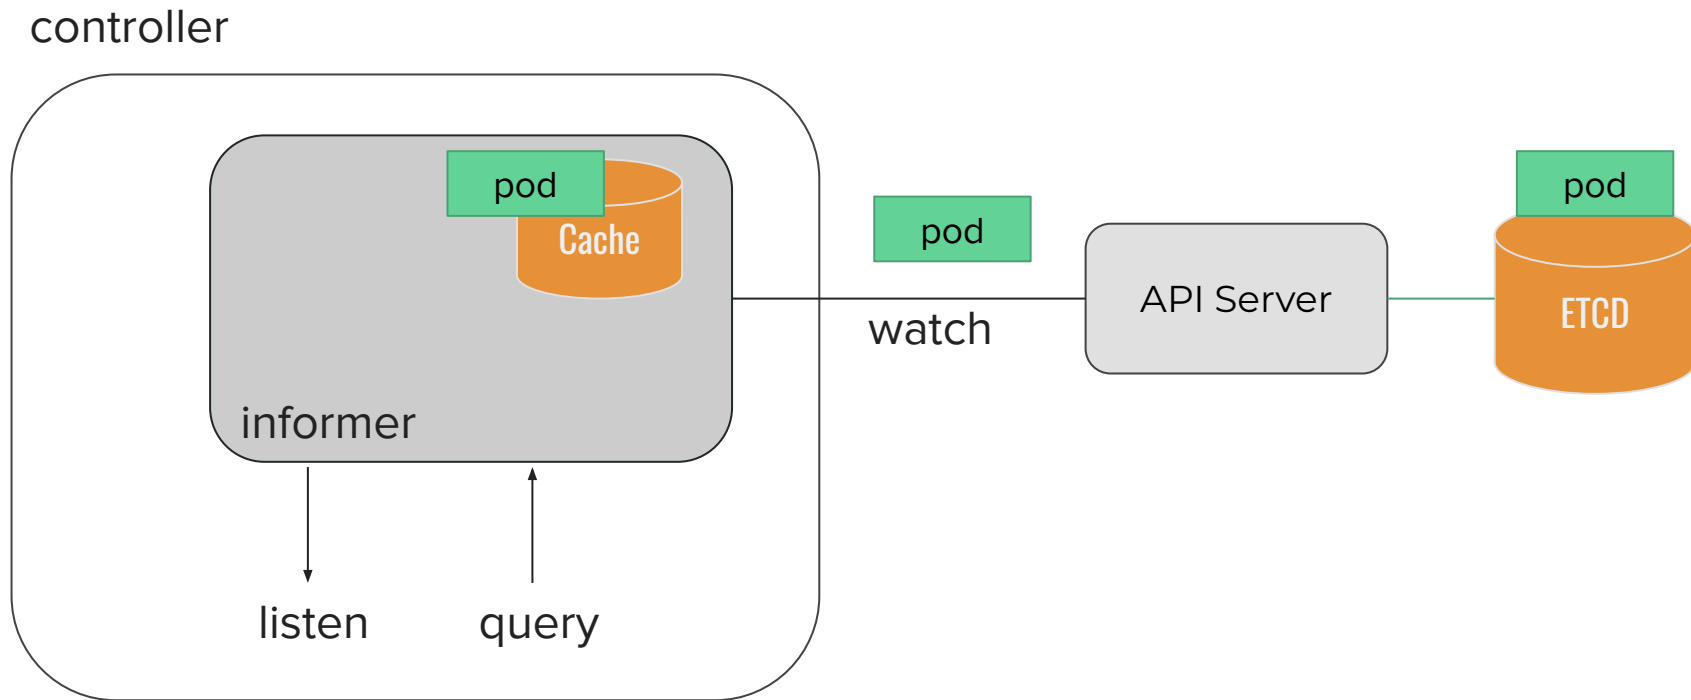




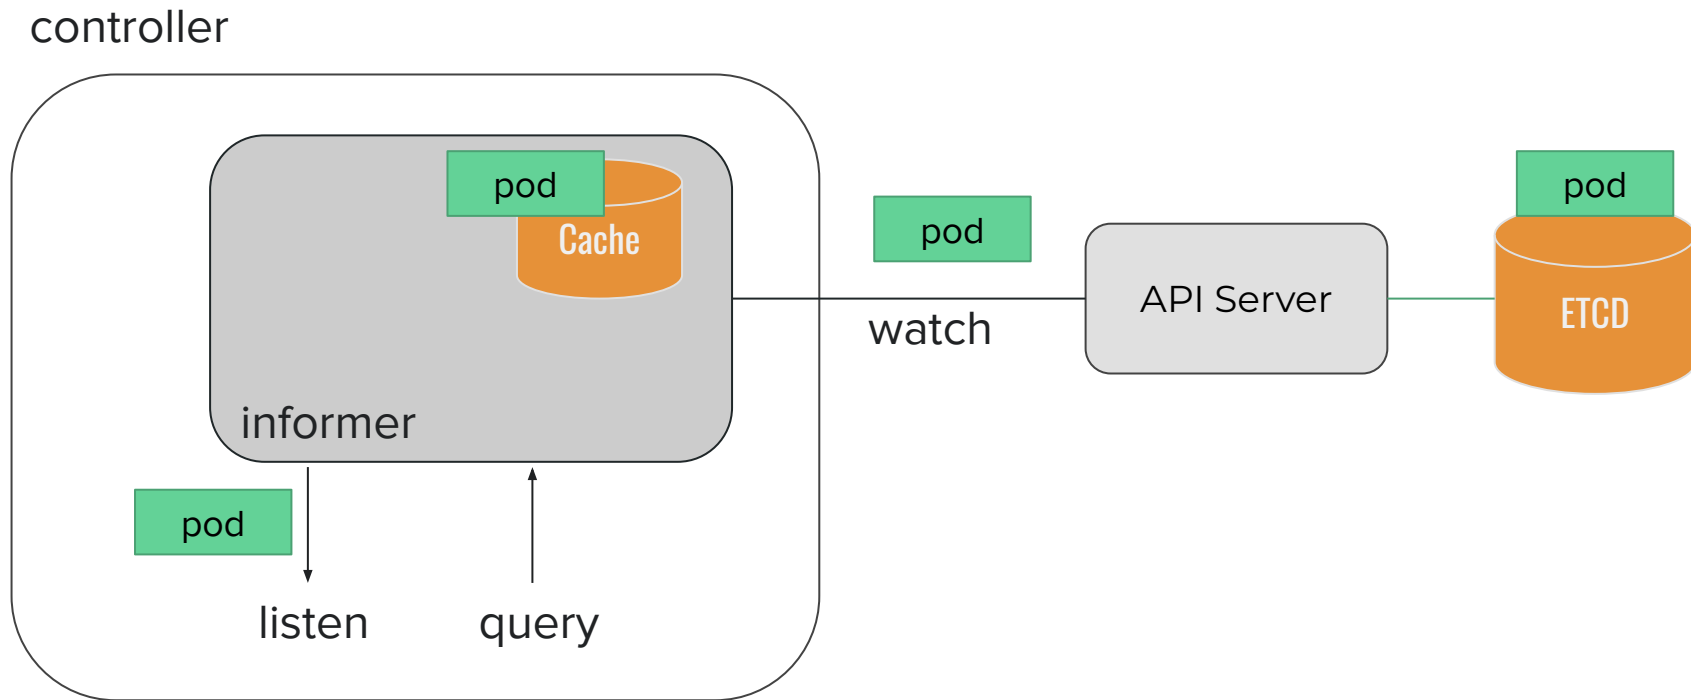
# Informers



# Informers



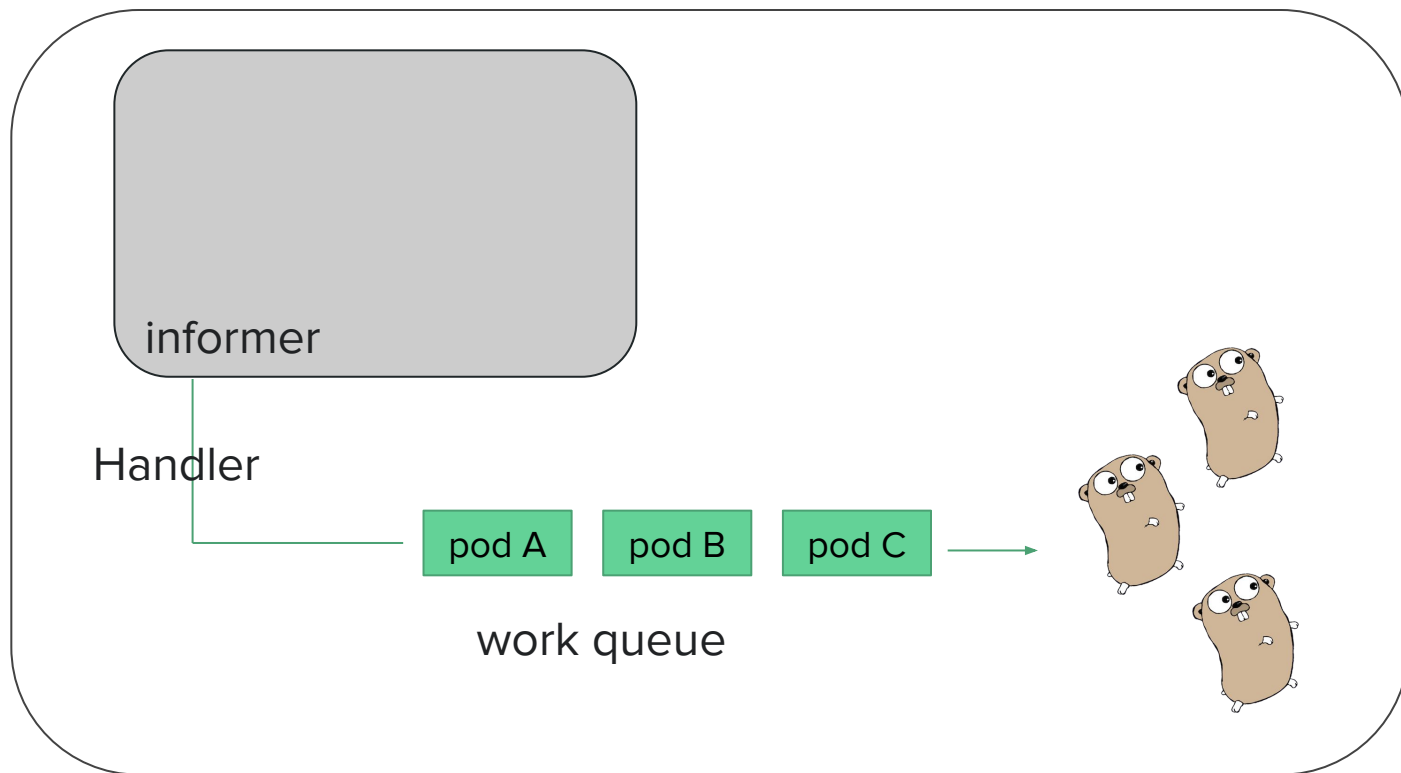
# Informers



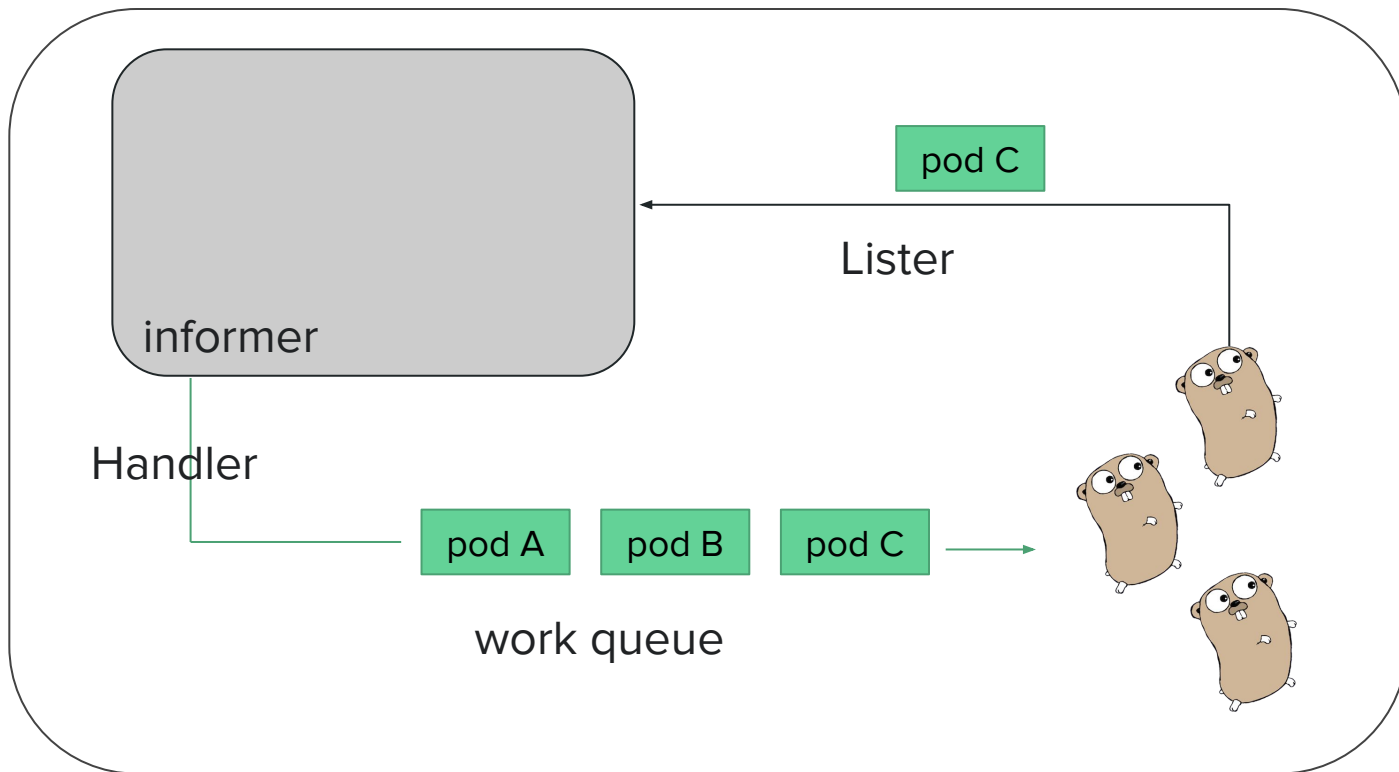
# Work queues

- Queue implementation
- Contains only the key of the object
- A given key occurs only once
- In case of failure the key can be pushed back to the queue

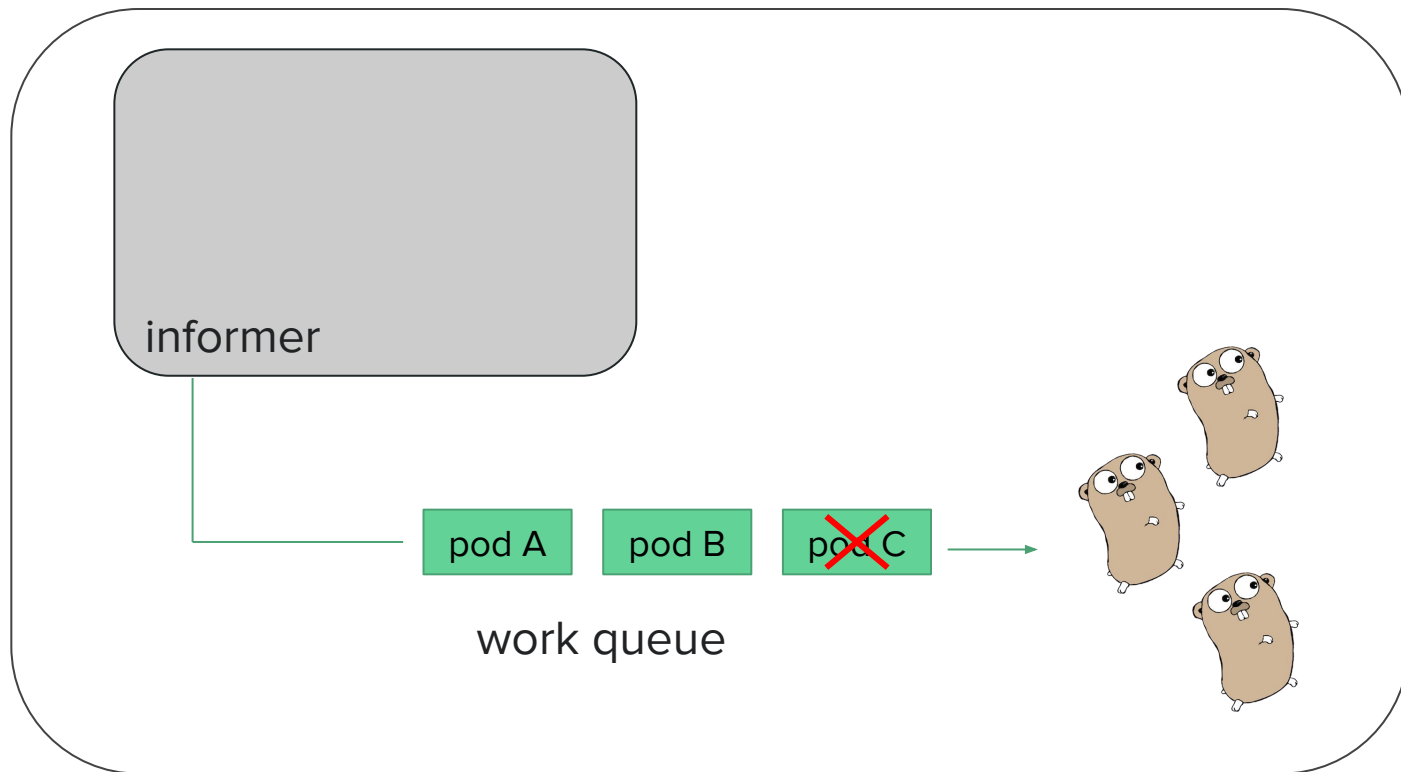
# Work queues




# Work queues



# Work queues



# Work queues





# Custom Resources Definitions

- Kubernetes types describing user defined kubernetes types
- Can be handled by controllers
- Represent our cloud-native API


# Custom Resources Definitions

```
apiVersion: apiextensions.k8s.io/v1
kind: CustomResourceDefinition
metadata:
  name: crontabs.stable.example.com
spec:
  group: stable.example.com
  versions:
    - name: v1
      served: true
      storage: true
```

```
schema:
  openAPIV3Schema:
    type: object
    properties:
      spec:
        type: object
        properties:
          cronSpec:
            type: string
          image:
            type: string
          replicas:
            type: integer
scope: Namespaced
names:
  plural: crontabs
  singular: crontab
  kind: CronTab
  shortNames:
    - ct
```

Spec is the desired state, status the actual state

```
apiVersion: example.io/v1alpha1
kind: Thermostat
metadata:
  name: high
  namespace: default
spec:
  temperature: 100
status:
  state: "on"
  temperature: 75
```




# Namespaces

---

# Namespaces


- Logical partitions across Kubernetes clusters
- Isolation of resources:
  - Pods
  - Configmaps
  - CRs
- Used for:
  - Multitenancy
  - Different environments

# Namespaces

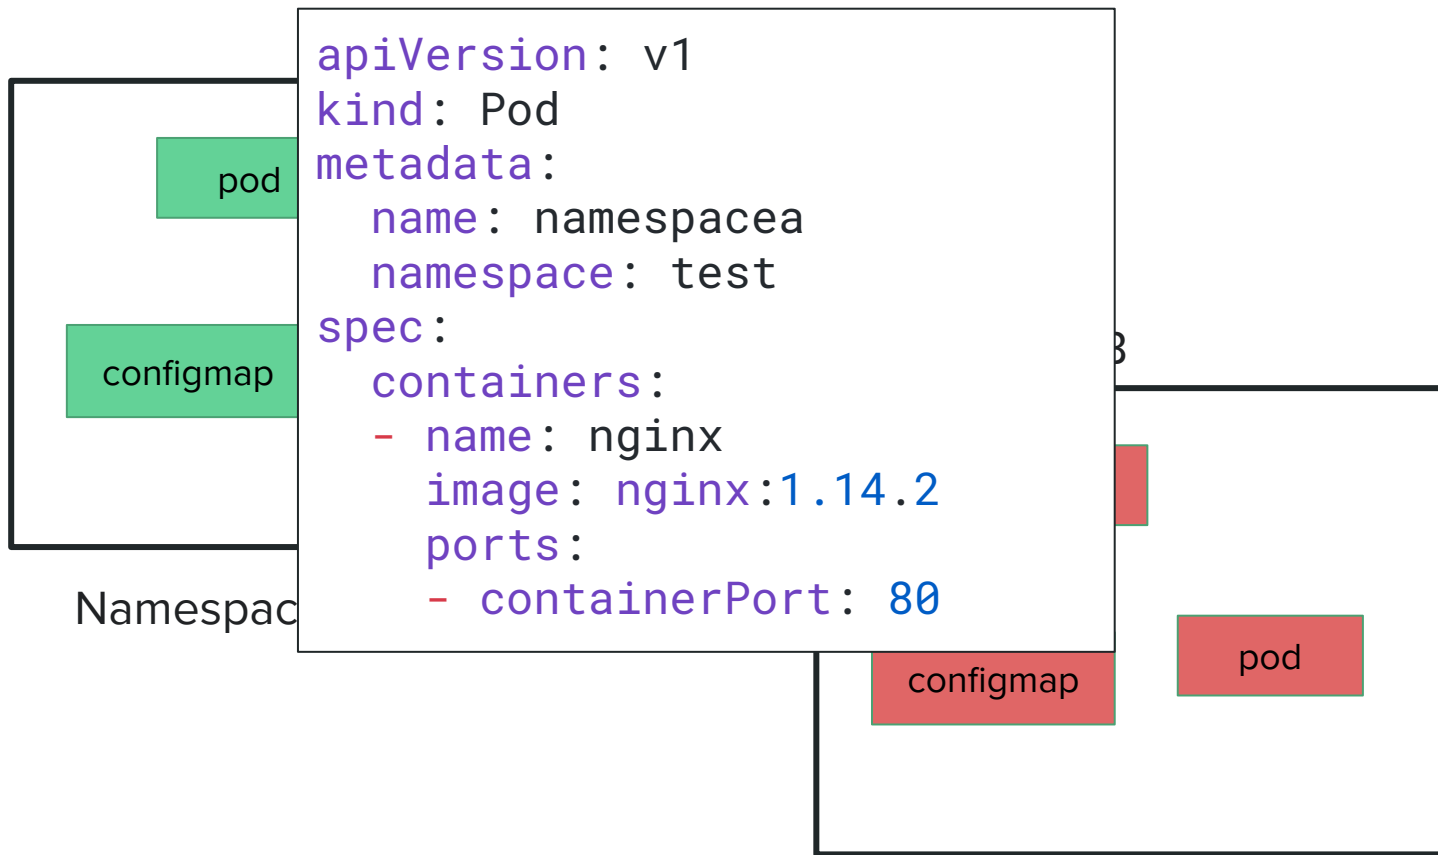


Namespace A

Namespace B



# Namespaces



# The permission model

---



RBAC?

Role

Based

Access

Control

## Role / Clusterrole

```
- apiGroups:  
  ["example.golab.io"]  
  resources:  
    - thermostats  
  verbs:  
    - get  
    - list  
    - watch
```

Role / Clusterrole

Subject  
(Service Account)

```
apiVersion: v1
kind: ServiceAccount
metadata:
  name: golabcontroller
```

Role / Clusterrole



```
graph TD; A[Role / Clusterrole] --> B[Role / Clusterrole Binding]; B --> C[Subject (Service Account)]
```

Role / Clusterrole  
Binding

Subject  
(Service Account)

```
apiVersion:
rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  name: golab:controller
roleRef:
  apiGroup:
rbac.authorization.k8s.io
  kind: ClusterRole
  name: golab:controller
subjects:
- kind: ServiceAccount
  name: golabcontroller
```

## Role vs ClusterRole

- Cluster role provides per cluster permissions
- Role provides per namespace permissions

# Too Many Yaml!

---

# Kubebuilder

- A framework for building Kubernetes APIs using Go.
- Simplifies creation of Custom Resource Definitions (CRDs) and controllers.
- Part of the official Kubernetes SIG API Machinery.
- Leverages Controller Runtime

[kubebuilder.io](https://kubebuilder.io)

# What Kubebuilder does

- Generates the project scaffolding
- Generates the yamls for the CRDs
- Generates the yamls for the permissions
- Generates the boilerplate code for the reconciliation loop

[kubebuilder.io](https://kubebuilder.io)



# Project generation

---

# Kubebuilder

```
kubebuilder init --domain golab.io --repo golab.io/kubedredger
```

# Kubebuilder

```
kubebuilder init --domain golab.io --repo golab.io/kubedredger
```

```
cmd  
config  
Dockerfile  
go.mod  
go.sum  
hack  
Makefile  
PROJECT  
README.md  
test
```

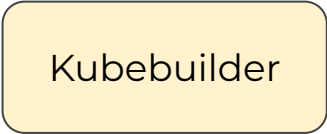
# Makefile rules

- make docker-build
- make build
- make lint
- make test
- make setup-e2e
- make e2e
- ....

API

---


# The workflow




Kubebuilder

```
kubebuilder create api --group workshop --version v1alpha1 \  
--kind Configuration
```

# The workflow



# The workflow




```
api
├── v1alpha1
│   ├── configuration_types.go
│   ├── groupversion_info.go
│   └── zz_generated.deepcopy.go
```

```
// ConfigurationSpec defines the desired state of Configuration
type ConfigurationSpec struct {
    // INSERT ADDITIONAL SPEC FIELDS - desired state of cluster
    Foo *string `json:"foo,omitempty"`
}
```




# The workflow



```
// ConfigurationSpec defines the desired state of Configuration


type ConfigurationSpec struct {
    // Content is the content to be written to the file
    Content string `json:"content"`
    // Create indicates whether to create the file if it does not exist
    Create bool `json:"create,omitempty"`
    // +optional
    Permission *uint32 `json:"permission,omitempty"`
}
```

# The workflow



make manifests

# The workflow



## make manifests

```
apiVersion: apiextensions.k8s.io/v1
kind: CustomResourceDefinition
metadata:
  annotations:
    controller-gen.kubebuilder.io/version: v0.18.0
  name: configurations.workshop.golab.io
spec:
  group: workshop.golab.io
  names:
    kind: Configuration
    listKind: ConfigurationList
    plural: configurations
    singular: configuration
  scope: Namespaced
  versions:
    - name: v1alpha1
      schema:
        openAPIV3Schema:
          description: Configuration is the Schema for the
            configurations API
```

# API Properties

---

# Kubebuilder

```
// Foo is a foo field
// +kubebuilder:validation:Minimum=1
// +kubebuilder:default:=32
// +optional
Foo *int32 `json:"aggregationLength,omitempty"`
```

# Kubebuilder

```
// Foo is a foo field
// +kubebuilder:validation:Minimum=1
// +kubebuilder:default:=32
// +optional
Foo *int32 `json:"aggregationLength,omitempty"`
```

# Kubebuilder

```
// Foo is a foo field
// +kubebuilder:validation:Minimum=1
// +kubebuilder:default:=32
// +optional
Foo *int32 `json:"aggregationLength,omitempty"`
```

# Kubebuilder

```
// Foo is a foo field
// +kubebuilder:validation:Minimum=1
// +kubebuilder:default:=32
// +optional
Foo *int32 `json:"aggregationLength,omitempty"`
```



# Kubebuilder

```
// Foo is a foo field
// +kubebuilder:validation:Minimum=1
// +kubebuilder:default:=32
// +optional
Foo *int32 `json:"aggregationLength,omitempty"`
```

# Print columns

```
//
+kubebuilder:printcolumn:name="Path",type="string",JSONPath=".spec.path",description=
"Path of the file"
//
+kubebuilder:printcolumn:name="Exists",type="boolean",JSONPath=".status.fileExists",d
escription="Tells if the file exists"
//
+kubebuilder:printcolumn:name="LastUpdate",type="date",JSONPath=".status.lastUpdated"
,description="Last update of the file"
```

NAMESPACE  
LASTUPDATE

NAME

PATH

EXISTS

k8s-example-kubedredger-system

test-config

/foo.txt

true

16s

# RBAC generation

---

# RBAC

```
//+kubebuilder:rbac:groups=workshop.golab.io,resources=configurations,verbs=get;list;watch;create;update;patch;delete
// +kubebuilder:rbac:groups=workshop.golab.io,resources=configurations/status,verbs=get;update;patch
// +kubebuilder:rbac:groups=workshop.golab.io,resources=configurations/finalizers,verbs=update
```

make manifests

# RBAC

```
//+kubebuilder:rbac:
h;delete
// +kubebuilder:rbac
// +kubebuilder:rbac
```

make manifest

```
apiVersion:
rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: manager-role
rules:
- apiGroups:
  - workshop.golab.io
  resources:
  - configurations
  verbs:
  - create
  - delete
  - get
  - list
  - patch
  - update
  - watch
```

atch;create;update;patc

;update;patch

s=update

# KIND: Running a cluster on your laptop

---

# KIND: Kubernetes in Docker

- a tool for running local Kubernetes clusters using Docker container “nodes”.
- primarily designed for testing Kubernetes itself, but may be used for local development or CI.
- you need: kind itself, [kubectI](#), docker (or podman)

[kind.sigs.k8s.io](https://kind.sigs.k8s.io)

# KIND: for this workshop...


<https://github.com/goworkshops/k8s-example-kubedredger>

- make help - generic usage
- **make dep-install-kubectrl**
- **make dep-install-kind**



## Workshop-specific helpers

<code>dep-install-kubectrl</code>	Download kubectrl locally if necessary, or reuse the system binary.
<code>dep-install-golangci-lint</code>	Download golangci-lint locally if necessary, or reuse the system binary
<code>dep-install-kind</code>	Download kind locally if necessary, or reuse the system binary
<code>system-lint</code>	Run golangci-lint linter using system binaries
<code>system-lint-fix</code>	Run golangci-lint linter and perform fixes using system binaries
<code>system-lint-config</code>	Verify golangci-lint linter configuration using system binaries
<code>kind-setup</code>	Set up a Kind cluster if it does not exist
<code>kind-teardown</code>	Tear down the Kind cluster created with kind-setup
<code>kind-load-image</code>	Load the image in the local cluster





## KIND: for this workshop... /2

- in the repo root
- setup the kind cluster: **make kind-setup**
- teardown once done: **make kind-teardown**
- Your cluster is ready!

# KIND: for this workshop...

```
$ make kind-setup
Kind cluster 'k8s-example-kubedredger-kind' already exists. Skipping creation.
$ make kind-teardown
Deleting cluster "k8s-example-kubedredger-kind" ...
Deleted nodes: ["k8s-example-kubedredger-kind-control-plane"]
$ make kind-setup
No kind clusters found.
Creating Kind cluster 'k8s-example-kubedredger-kind'...
Creating cluster "k8s-example-kubedredger-kind" ...
  ✓ Ensuring node image (kindest/node:v1.34.0) 🖼️
  ✓ Preparing nodes 📦
  ✓ Writing configuration 📄
  ✓ Starting control-plane 🚦
  ✓ Installing CNI 🖱️
  ✓ Installing StorageClass 💾
Set kubectl context to "kind-k8s-example-kubedredger-kind"
You can now use your cluster with:

kubectl cluster-info --context kind-k8s-example-kubedredger-kind

Thanks for using kind! 😊
$ make kind-teardown
Deleting cluster "k8s-example-kubedredger-kind" ...
Deleted nodes: ["k8s-example-kubedredger-kind-control-plane"]
$
```

# Kubectl: accessing a kubernetes cluster

What is it: a command line tool for communicating with a Kubernetes cluster's control plane, using the Kubernetes API. Usage examples.

If you don't have it already:

- make kubectl # from the repo root

# Kubectrl: example 1: checking versions

```
$ kubectl version  
Client Version: v1.34.1  
Kustomize Version: v5.7.1  
Server Version: v1.34.0
```

## Kubect1: example 2: getting “nodes”

```
$ kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
k8s-example-kubedredger-kind-control-plane	Ready	control-plane	2m34s	v1.34.0

need more? please check `hack/kind*.yaml`

# Kubectl: example 3: getting all pods

```
$ kubectl get pods -A
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
kube-system	coredns-66bc5c9577-b5fp9	1/1	Running	0	2m37s
kube-system	coredns-66bc5c9577-q6zk9	1/1	Running	0	2m37s
kube-system	etcd-k8s-example-kubedredger-kind-control-plane	1/1	Running	0	2m44s
kube-system	kindnet-ksf8f	1/1	Running	0	2m37s
kube-system	kube-apiserver-k8s-example-kubedredger-kind-control-plane	1/1	Running	0	2m44s
kube-system	kube-controller-manager-k8s-example-kubedredger-kind-control-plane	1/1	Running	0	2m44s
kube-system	kube-proxy-vpgbr	1/1	Running	0	2m37s
kube-system	kube-scheduler-k8s-example-kubedredger-kind-control-plane	1/1	Running	0	2m44s
local-path-storage	local-path-provisioner-7b8c8ddb6-8tkd4	1/1	Running	0	2m37s

```
$ kubectl get pods -n local-path-storage
```

NAME	READY	STATUS	RESTARTS	AGE
local-path-provisioner-7b8c8ddb6-8tkd4	1/1	Running	0	19m

# Kubectrl: example 4: getting a pod YAML

```
$ kubectl get pod -o yaml -n kube-system kindnet-ksf8f
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  creationTimestamp: "2025-09-28T20:17:02Z"
```

```
  generateName: kindnet-
```

```
  generation: 1
```

```
  labels:
```

```
    app: kindnet
```

```
    controller-revision-hash: 57c957c676
```

```
    k8s-app: kindnet
```

```
    pod-template-generation: "1"
```

```
  tier: node
```

```
  name: kindnet-ksf8f
```

```
...
```

# Kubectctl: example 5: getting container logs

```
$ kubectl logs -n kube-system kindnet-ksf8f
I0928 20:17:04.161842      1 main.go:390] probe TCP address
k8s-example-kubedredger-kind-control-plane:6443
I0928 20:17:04.163231      1 main.go:109] connected to apiserver:
https://k8s-example-kubedredger-kind-control-plane:6443
I0928 20:17:04.163386      1 main.go:139] hostIP = 172.18.0.2
podIP = 172.18.0.2
I0928 20:17:04.163533      1 main.go:148] setting mtu 1500 for CNI
I0928 20:17:04.163545      1 main.go:178] kindnetd IP family: "ipv4"
I0928 20:17:04.163555      1 main.go:182] noMask IPv4 subnets: [10.244.0.0/16]
...
```



Kubect! (many) more examples

[kubect! quick reference](#)

# Running in a cluster

---

# How to run my container?

1. create container image
2. upload container image to a registry
3. send manifests to the cluster referring the image
4. success!

# How to run my container? for this workshop

1. make container-image
2. make kind-load-image
3. make deploy
4. success!

# Make container-image

Create the container image based on the Dockerfile

calls docker build with the correct parameters

# Make kind-load-image

regenerate partial manifests from annotations (implied dependency)

create manifests from the template

sends the manifests to the (kind) cluster

# Make deploy

Load the controller container image in the kind cluster directly in its data store

- avoids the need to push into a container registry
- requires **imagePullPolicy: ifNotPresent** in the manifests referring the image
- otherwise kubernetes will try to use a registry

# How to run my container? for this workshop

Optional but recommended activity: try to look the targets in the Makefile and look at the commands they run

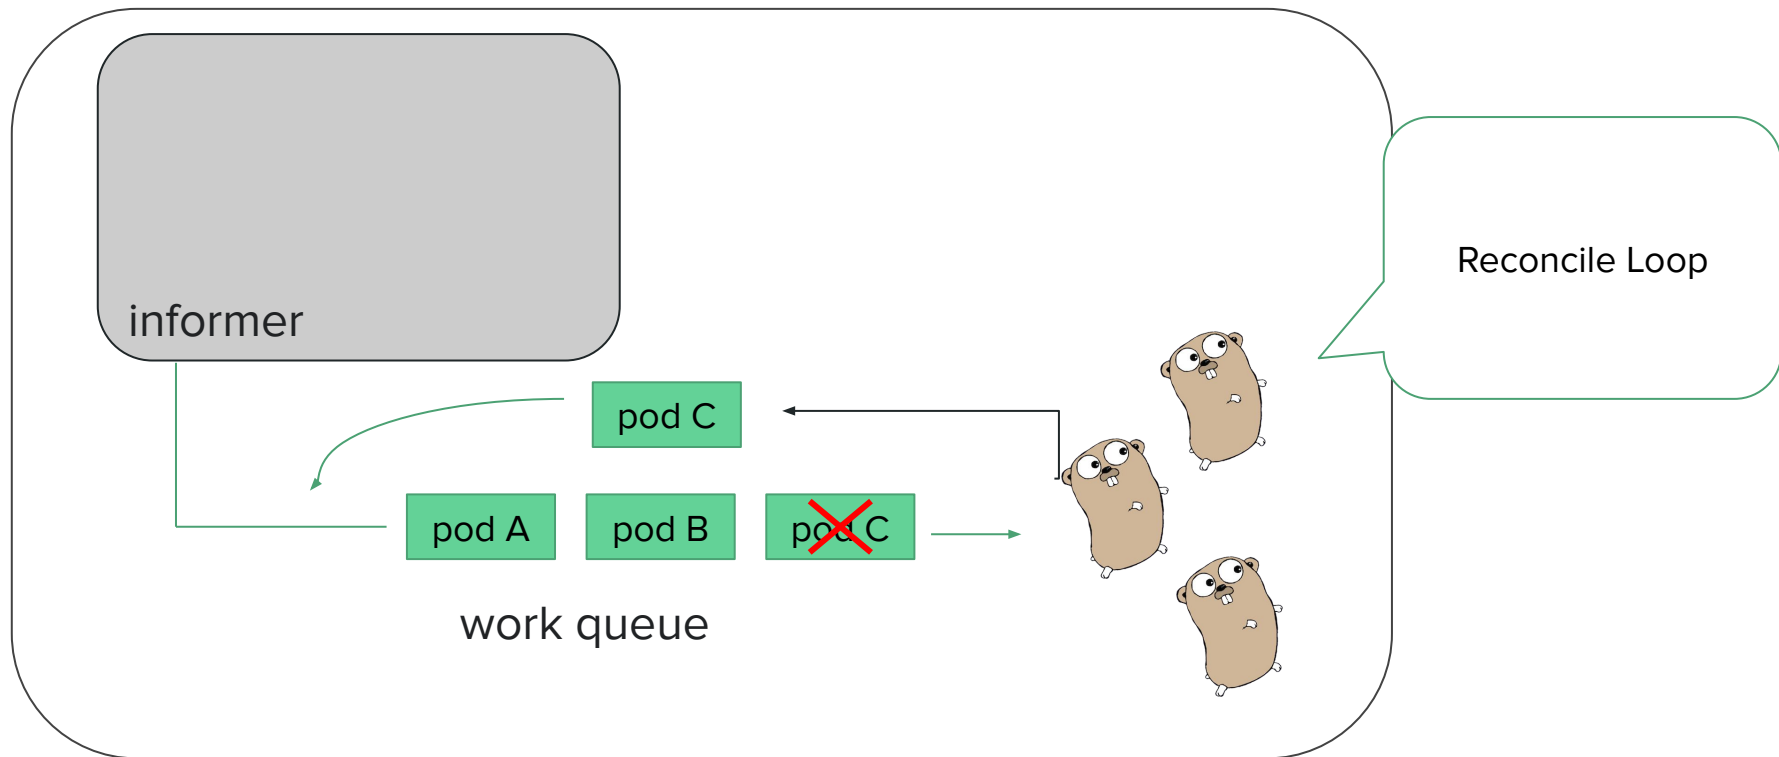
1. make container-image
2. make kind-load-image
3. make deploy




# Writing the reconcile loop

---

# What's a controller?



# What's a controller?



Kubebuilder uses  
controller-runtime

Controller runtime is a convenient wrapper of clients, caches, informers, watches

## Building Blocks

- Reconcile loop
- Cache / Client to access resources

# Automatic scaffolding generation

```
kubebuilder create api --group workshop --version v1alpha1 --kind
Configuration
INFO Create Resource [y/n]
y
INFO Create Controller [y/n]
y
INFO Writing kustomize manifests for you to edit...
INFO Writing scaffold for you to edit...
INFO api/v1alpha1/configuration_types.go
INFO api/v1alpha1/groupversion_info.go
INFO internal/controller/suite_test.go
INFO internal/controller/configuration_controller.go
INFO internal/controller/configuration_controller_test.go
```

# What we get

```
type ConfigurationReconciler struct {  
    client.Client  
    Scheme *runtime.Scheme  
}
```

# What we get

```
type ConfigurationReconciler struct {  
    client.Client  
    Scheme *runtime.Scheme  
}
```

Our controller type

- Implements the reconcile loop
- Hosts fields needed to implement the business logic and to tie Kubernetes to the real world



```
//
+kubebuilder:rbac:groups=workshop.golab.io,resources=configurations,verbs=get;list;watch;create;update;patch;delete
//
+kubebuilder:rbac:groups=workshop.golab.io,resources=configurations/status,verbs=get;update;patch
//
+kubebuilder:rbac:groups=workshop.golab.io,resources=configurations/finalizers,verbs=update
```

```
func (r *ConfigurationReconciler) Reconcile(ctx context.Context, req ctrl.Request) (ctrl.Result, error) {
    _ = logf.FromContext(ctx)

    // TODO(user): your logic here

    return ctrl.Result{}, nil
}
```

```
//  
+kubebuilder:rbac:groups=workshop.golangci.io,resources=workshop.golangci.io,verbs=get;list;watch;create;update;patch;delete  
//  
+kubebuilder:rbac:groups=workshop.golangci.io,resources=workshop.golangci.io,verbs=get;update;patch  
//  
+kubebuilder:rbac:groups=workshop.golangci.io,resources=workshop.golangci.io,verbs=update
```

The reconcile loop. Gets triggered every time a resource we are monitoring is changed

```
func (r *ConfigurationReconciler) Reconcile(ctx context.Context, req ctrl.Request) (ctrl.Result, error) {  
    _ = logf.FromContext(ctx)  
  
    // TODO(user): your logic here  
  
    return ctrl.Result{}, nil  
}
```

```
// SetupWithManager sets up the controller with the Manager.  
func (r *ConfigurationReconciler) SetupWithManager(mgr ctrl.Manager)  
error {  
    return ctrl.NewControllerManagedBy(mgr).  
        For(&workshopvalpha.Configuration{}).  
        Named("configuration").  
        Complete(r)  
}
```

Wiring to the controller manager.  
We declare the objects we are interested into

```
// SetupWithManager sets up the controller with the Manager.  
func (r *ConfigurationReconciler) SetupWithManager(mgr ctrl.Manager)  
error {  
    return ctrl.NewControllerManagedBy(mgr).  
        For(&workshopvalpha.Configuration{}).  
        Named("configuration").  
        Complete(r)  
}
```

Let's focus on Reconcile

---


```
func (r *ConfigurationReconciler) Reconcile(ctx context.Context,  
                                             req ctrl.Request) (ctrl.Result, error) {  
    lh := logf.FromContext(ctx)  
  
    conf := workshopv1alpha1.Configuration{}  
    err := r.Get(ctx, req.NamespacedName, &conf)  
  
    if apierrors.IsNotFound(err) {  
        // handle deletion  
    }  
}
```

```
type NamespacedName struct {  
    Namespace string  
    Name      string  
}
```

```
func (r *ConfigurationReconciler) Reconcile(ctx context.Context,  
                                             req ctrl.Request) (ctrl.Result, error) {  
    lh := logf.FromContext(ctx)  
  
    conf := workshopv1alpha1.Configuration{}  
    err := r.Get(ctx, req.NamespacedName, &conf)  
  
    if apierrors.IsNotFound(err) {  
        // handle deletion  
    }  
}
```

Req contains the key of the object, not the object itself

- To access the object, it must be fetched using the client
- No object means the object was cancelled
- No old values and new values!





# The return values

```
// If the returned error is non-nil, the Result is ignored and
// the request will be requeued using exponential backoff.
// The only exception is if the error is a TerminalError
// in which case no requeuing happens.
//
// If the error is nil and the returned Result has a non-zero
// result.RequeueAfter, the request will be requeued after
// the specified duration.
```

```
Reconcile(context.Context, request) (Result, error)
```

# The return values

```
// If the returned error is non-nil, the Result is ignored and  
// the request will be requeued using exponential backoff.  
// The only exception is if the error is a TerminalError  
// in which case no requeuing happens.  
//  
// If the error is nil and the returned Result has a non-zero  
// result.RequeueAfter, the request will be requeued after  
// the specified duration.
```

```
Reconcile(context.Context, request) (Result, error)
```

# The return values

```
// If the returned error is non-nil, the Result is ignored and
// the request will be requeued using exponential backoff.
// The only exception is if the error is a TerminalError
// in which case no requeuing happens.
//
// If the error is nil and the returned Result has a non-zero
// result.RequeueAfter, the request will be requeued after
// the specified duration.
```

```
Reconcile(context.Context, request) (Result, error)
```

# The return values

```
// If the returned error is non-nil, the Result is ignored and
// the request will be requeued using exponential backoff.
// The only exception is if the error is a TerminalError
// in which case no requeuing happens.
//
// If the error is nil and the returned Result has a non-zero
// result.RequeueAfter, the request will be requeued after
// the specified duration.
```

```
Reconcile(context.Context, request) (Result, error)
```


# Anatomy of a good controller

---

# Anatomy of a good controller

- business logic has a clear boundary, with a its own API
- the reconciliation logic is IDEMPOTENT
- any transient error must be sent back to the queue
- any non transient error must be logged and Reconcile must return success

```
type ConfigurationReconciler struct {  
    client.Client  
    Scheme *runtime.Scheme  
    ConfMgr *configfile.Manager  
}
```



The business logic we want to implement

```
err := r.Get(ctx, req.NamespacedName, &conf)
if apierrors.IsNotFound(err) {
    if err := r.ConfMgr.Delete(); err != nil {
        return ctrl.Result{}, fmt.Errorf("failed to delete the
configuration: %w", err)
    }
    return ctrl.Result{}, nil
}

if err != nil {
    // Error reading the object - requeue the request.
    return ctrl.Result{}, err
}
```



```
err := r.Get(ctx, req.NamespacedName, &conf)
if apierrors.IsNotFound(err) {
    if err := r.ConfMgr.Delete(); err != nil {
        return ctrl.Result{},
configuration: %w", err)
    }
    return ctrl.Result{}, nil
}

if err != nil {
    // Error reading the object - requeue the request.
    return ctrl.Result{}, err
}
```

Handle the deletion

the

```
err := r.Get(ctx, req.NamespacedName, &conf)
if apierrors.IsNotFound(err) {
    if err := r.ConfMgr.Delete(); err != nil {
        return ctrl.Result{}, fmt.Errorf("failed to delete
configuration: %w", err)
    }
    return ctrl.Result{}, nil
}

if err != nil {
    // Error reading the object - requeue the request.
    return ctrl.Result{}, err
}
```

Requeue the request

the

```
configurationRequest := configurationRequestFromSpec(conf.Spec)

err = r.ConfMgr.HandleSync(lh, configurationRequest)
if errors.As(err, &configfile.NonRecoverableError{}) {
    lh.Error(err, "Non-recoverable error handling
configuration")
    return ctrl.Result{}, nil
}
```

Translation from the kubernetes  
API to internal

```
configurationRequest := configurationRequestFromSpec(conf.Spec)

err = r.ConfMgr.HandleSync(lh, configurationRequest)
if errors.As(err, &configfile.NonRecoverableError{}) {
    lh.Error(err, "Non-recoverable error handling
configuration")
    return ctrl.Result{}, nil
}
```

Business logic with its own API  
Must be idempotent

```
configurationRequest := confi  
  
err = r.ConfMgr.HandleSync(lh, configurationRequest)  
if errors.As(err, &configfile.NonRecoverableError{}) {  
    lh.Error(err, "Non-recoverable error handling  
configuration")  
    return ctrl.Result{}, nil  
}
```

# Notes

- the cache gets warmed up before reconcile is called
- when we get an object using Get or List, we access the local cache and not the latest data on the api server
- because of how the informer works, the cache is updated before the handler is called

# Takeaways

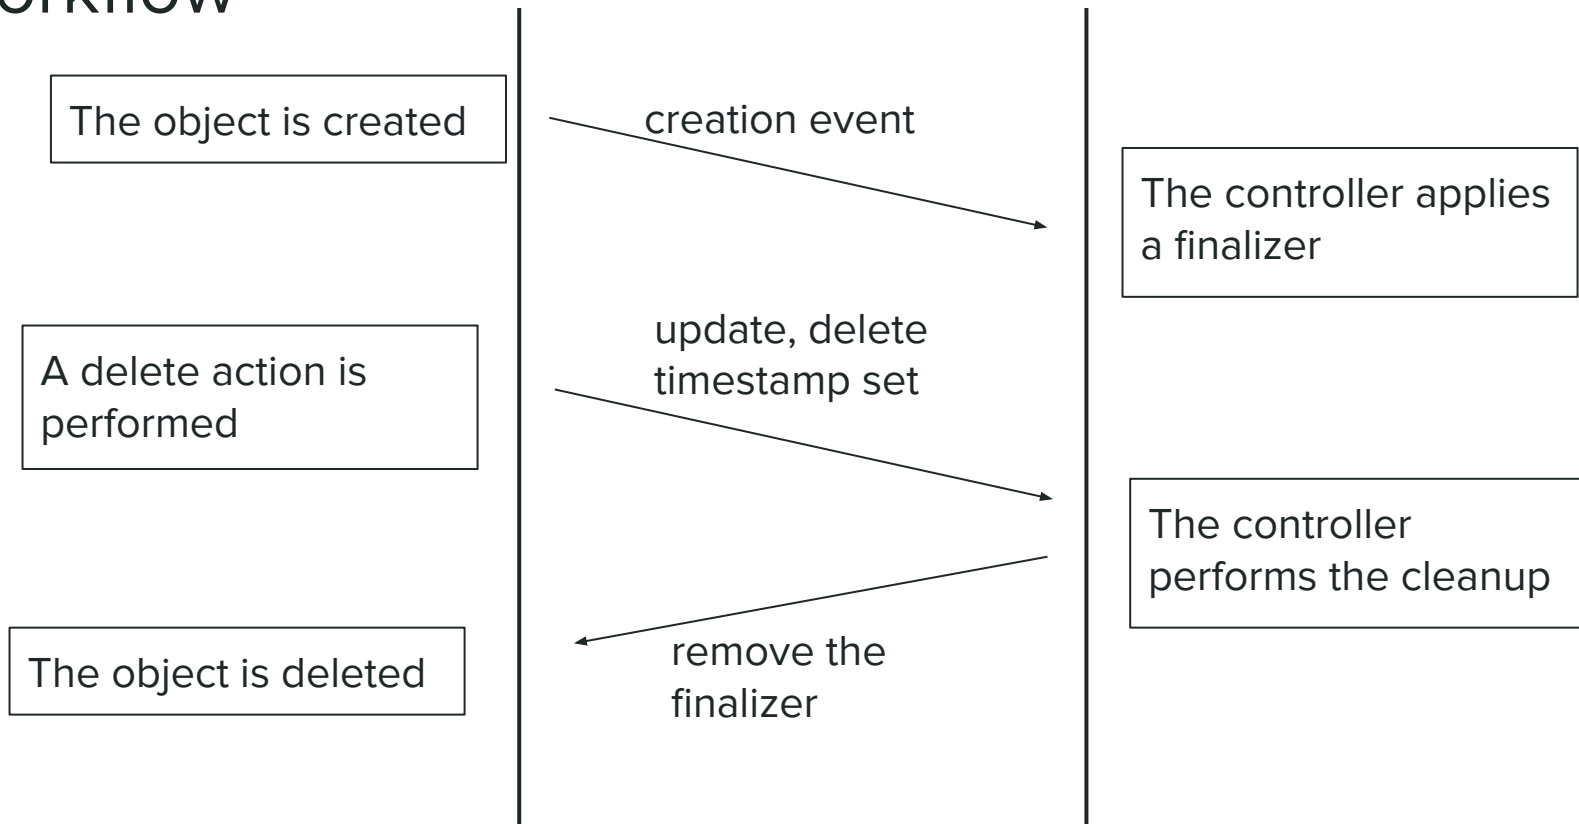
- we don't know how many times our reconcile is getting called with the same data
- non-k8s business logic allows us to test it in isolation
- the business logic must be idempotent
- the reward is a bullet proof controller

# Finalizers

---



# The workflow



## Finalizers are pre-delete hooks

- An object does not get removed as long as it has a finalizer applied
- This gives time to perform housekeeping asynchronously
- Also, it allow us to use the whole object before it gets deleted

# Finalizers require specific permissions

```
//+kubebuilder:rbac:groups=workshop.golab.io,resources=configurations,verbs=get;list;watch;create;update;patch;delete
// +kubebuilder:rbac:groups=workshop.golab.io,resources=configurations/status,verbs=get;update;patch
// +kubebuilder:rbac:groups=workshop.golab.io,resources=configurations/finalizers,verbs=update
```

# In the reconcile loop

```
if !conf.DeletionTimestamp.IsZero() {
    if controllerutil.ContainsFinalizer(conf, Finalizer) {

        // PERFORM DELETE
        controllerutil.RemoveFinalizer(conf, Finalizer)
        err = r.Update(ctx, conf)
        return ctrl.Result{}, err
    }
    return ctrl.Result{}, nil
}

if !controllerutil.ContainsFinalizer(conf, Finalizer) {
    controllerutil.AddFinalizer(conf, Finalizer)
    err = r.Update(ctx, conf)
    if err != nil {
        return ctrl.Result{}, err
    }
}

// RECONCILE LOGIC
```

# In the reconcile loop

It's a delete!

```
if !conf.DeletionTimestamp.IsZero() {
    if controllerutil.ContainsFinalizer(conf, Finalizer) {

        // PERFORM DELETE
        controllerutil.RemoveFinalizer(conf, Finalizer)
        err = r.Update(ctx, conf)
        return ctrl.Result{}, err
    }
    return ctrl.Result{}, nil
}

if !controllerutil.ContainsFinalizer(conf, Finalizer) {
    controllerutil.AddFinalizer(conf, Finalizer)
    err = r.Update(ctx, conf)
    if err != nil {
        return ctrl.Result{}, err
    }
}

// RECONCILE LOGIC
```

# In the reconcile loop

If the finalizer is set, we need to perform the delete (or the delete is in progress)

```
if !conf.DeletionTimestamp.IsZero() {
    if controllerutil.ContainsFinalizer(conf, Finalizer) {

        // PERFORM DELETE
        controllerutil.RemoveFinalizer(conf, Finalizer)
        err = r.Update(ctx, conf)
        return ctrl.Result{}, err
    }
    return ctrl.Result{}, nil
}

if !controllerutil.ContainsFinalizer(conf, Finalizer) {
    controllerutil.AddFinalizer(conf, Finalizer)
    err = r.Update(ctx, conf)
    if err != nil {
        return ctrl.Result{}, err
    }
}

// RECONCILE LOGIC
```

# In the reconcile loop

```
if !conf.DeletionTimestamp.IsZero() {  
    if controllerutil.ContainsFinalizer(conf, Finalizer) {  
  
        // PERFORM DELETE  
        controllerutil.RemoveFinalize  
        err = r.Update(ctx, conf)  
        return ctrl.Result{}, err  
    }  
    return ctrl.Result{}, nil  
}
```

If not, the object is just being deleted, we can exit

```
if !controllerutil.ContainsFinalizer(conf, Finalizer) {  
    controllerutil.AddFinalizer(conf, Finalizer)  
    err = r.Update(ctx, conf)  
    if err != nil {  
        return ctrl.Result{}, err  
    }  
}
```

```
// RECONCILE LOGIC
```

# In the reconcile loop

```
if !conf.DeletionTimestamp.IsZero() {  
    if controllerutil.ContainsFinalizer(conf, Finalizer) {  
  
        // PERFORM DELETE  
        controllerutil.RemoveFinalizer(conf, Finalizer)  
        err = r.Update(ctx, conf)  
        return ctrl.Result{}, err  
    }  
    return ctrl.Result{}, nil  
}  
  
if !controllerutil.ContainsFinalizer(conf, Finalizer) {  
    controllerutil.AddFinalizer(conf, Finalizer)  
    err = r.Update(ctx, conf)  
    if err != nil {  
        return ctrl.Result{}, err  
    }  
}
```

Add / update. We want to set the finalizer so we control the real deletion of the object

```
// RECONCILE LOGIC
```



# In the reconcile loop

```
if !conf.DeletionTimestamp.IsZero() {  
    if controllerutil.ContainsFinalizer(conf, Finalizer) {  
  
        // PERFORM DELETE  
        controllerutil.RemoveFinalizer(conf, Finalizer)  
        err = r.Update(ctx, conf)  
        return ctrl.Result{}, err  
    }  
    return ctrl.Result{}, nil  
}  
  
if !controllerutil.ContainsFinalizer(conf, Finalizer) {  
    controllerutil.AddFinalizer(conf, Finalizer)  
    err = r.Update(ctx, conf)  
    if err != nil {  
        return ctrl.Result{Requeue: true}, err  
    }  
}  
  
// RECONCILE LOGIC
```

Regular reconciliation logic

# Exposing the status

---

## Status Part of the spec

```
type Pod struct {  
    metav1.TypeMeta  
    metav1.ObjectMeta  
    Spec    PodSpec  
    Status PodStatus  
}
```

```
type MyObject struct {  
    metav1.TypeMeta  
    metav1.ObjectMeta  
    Spec    MyObjectSpec  
    Status MyObjectStatus  
}
```


# Provide feedback to the users

```
kubectl get pods
```


NAME	READY	STATUS	RESTARTS	AGE
pod0	1 / 1	Running	0	1h
pod1	0 / 1	CrashLoopBackOff	8	16m

The status is machine friendly


# Feed other controllers



# Feed other controllers



# Status only objects



use case: enumerate resource and expose them  
using a kubernetes-native object

# The status is just another part of the CRD

```
// ConfigurationStatus defines the observed state of Configuration.
type ConfigurationStatus struct {
    // LastUpdated is the last time the configuration was updated
    LastUpdated metav1.Time `json:"lastUpdated"`

    // Content is the current content of the file at the specified path
    Content string `json:"content,omitempty"`

    // FileExists indicates whether the file exists at the specified path
    FileExists bool `json:"fileExists,omitempty"`
    // The status of each condition is one of True, False, or Unknown.
    // +listType=map
    // +listMapKey=type
    // +optional
    Conditions []metav1.Condition `json:"conditions,omitempty"`
}
```



# The status is just another part of the CRD

```
// ConfigurationStatus defines the observed state of Configuration.
type ConfigurationStatus struct {
    ...

    // The status of each condition is one of True, False, or Unknown.
    // +listType=map
    // +listMapKey=type
    // +optional
    Conditions []metav1.Condition `json:"conditions,omitempty"`
}
```



What conditions are?

It's a convention

It's a very good and very helpful convention we should strive to adhere to

## Why conditions?

The definition of status is object-dependent

Conditions are object-independent.

Conditions provide standard, generic, machine readable status information. Great for interoperability.

# Conditions

Each condition has:

- type (e.g., Available, Progressing, Degraded)
- status (True, False, Unknown)
- last transition time
- reason (for machines) & message (for humans)

# The status requires a different permission

```
//
```

```
+kubebuilder:rbac:groups=workshop.golab.io,resources=configurations/  
status,verbs=get;update;patch
```

- **apiGroups:**
  - workshop.golab.io
- **resources:**
  - configurations/status
- **verbs:**
  - get
  - patch
  - update

# A simple implementation

```
oldStatus := conf.Status.DeepCopy()

// use the conf object

confManagerStatus := r.ConfMgr.Status()
conf.Status = statusFromConfManagerStatus(conf.Spec, confManagerStatus, err)

if !statusesAreEqual(oldStatus, &conf.Status) {
    updErr := r.Client.Status().Update(ctx, &conf)
    if updErr != nil {
        lh.Error(updErr, "Failed to update configuration status")
        return ctrl.Result{}, fmt.Errorf("could not update status for object")
    }
}
```

# In our reconcile loop

Note the different call

```
oldStatus := conf.Status.DeepCopy()

// use the conf object

confManagerStatus := r.ConfMgr.Status()
conf.Status = statusFromConfManagerStatus(conf.Spec, confManagerStatus, err)

if !statusesAreEqual(oldStatus, &conf.Status) {
    updErr := r.Client.Status().Update(ctx, &conf)
    if updErr != nil {
        lh.Error(updErr, "Failed to update configuration status")
        return ctrl.Result{}, fmt.Errorf("could not update status for object")
    }
}
```

# In our reconcile loop

```
oldStatus := conf.Status.DeepCopy()

// use the conf object

confManagerStatus := r.ConfMgr.Status()
conf.Status = statusFromConfManagerStatus(conf.Spec, confManagerStatus, err)

if !statusesAreEqual(oldStatus, &conf.Status) {
    updErr := r.Client.Status().Update(ctx, &conf)
    if updErr != nil {
        lh.Error(updErr, "Failed to update configuration status")
        return ctrl.Result{}, fmt.Errorf("could not update status for object")
    }
}
```

If the update fails, we retry



# Testing a controller

---


# Makefile rules (reprised)

- make docker-build
- make build
- make lint
- **make test**
- **make test-e2e**
- ....


# Make test

- runs the regular unit test we know and love
  - the example project intentionally only uses the stdlib testing package
- also run the controller (semi)integration tests
- e2e (end to end) tests are run with an explicit target
- running a subset of tests:
  - `go test ./internal/...`


# The testing pyramid



# The testing pyramid



# The testing pyramid: fast vs slow



## unit tests: faking a go client

- The client-go package exposes a fake client implementation
- **ALMOST** native implementation look and feel
  - no server side -> no index support

```
import "sigs.k8s.io/controller-runtime/pkg/client/fake"
```

```
fake.NewClientBuilder().WithScheme(scheme.Scheme).Build()
```

## unit tests: faking a go client /2

You can create objects as usual, they will end up in the fake client store

You can pre-load objects in the store (skip explicit Create in your test)

```
import "k8s.io/apimachinery/pkg/runtime"
```

```
var objects []runtime.Object
```

```
fake.NewClientBuilder().WithScheme(scheme.Scheme).WithRuntimeObjects(objects...)
```



unit tests: faking a go client with indexer

Indexer support. E.g.

```
sel, err := fields.ParseSelector("spec.nodeName=" + nodeName)
...
podList := &corev1.PodList{}
cli.List(ctx, podList, &client.ListOptions{FieldSelector: sel})
```

## unit tests: faking a go client with indexer

```
func podNodeIndexer(rawObj client.Object) []string {  
    obj, ok := rawObj.(*corev1.Pod)  
    if !ok {  
        return []string{}  
    }  
    return []string{obj.Spec.NodeName}  
}
```

[reference](#) (note original issue deleted)

unit tests: faking a go client with indexer

```
fake.NewClientBuilder().  
    WithScheme(scheme.Scheme).  
    WithIndex(&corev1.Pod{}, "spec.nodeName", podNodeIndexer).  
    Build()
```

## putting all together: envtest

envtest is a helper package which manages instances of etcd and the Kubernetes API server, creating a micro-cluster-like without kubelets or any other controllers.

Bare bones to test a controller, without fakes

Scaffolded by kubebuilder

Use the regular standard client

## putting all together: envtest /2

```
testEnv = &envtest.Environment{
    CRDDirectoryPaths: []string{filepath.Join("../", "../", "config", "crd", "bases")},
    ErrorIfCRDPathMissing: true,
}

cfg, err = testEnv.Start()    // cfg is defined globally

k8sClient, err = client.New(cfg, client.Options{Scheme: scheme.Scheme})

// your test code here but beware: no other controller is running!

err = testEnv.Stop()
```

# Integration tests and beyond: Ginkgo and gomega

Ginkgo is a powerful testing framework for go

Gomega is a library which adds matching capabilities to ginkgo

(BeTrue, IsNil...)

Ginkgo and Gomega together provide a Domain-Specific Language (DSL) to write tests in go

# The ginkgo use case: asynchronous tests

```
Eventually(X).WithTimeout(T).WithPolling(P).WithContext(ctx).Should(MATCHER)
```

Checks an assertion passes *eventually*

- tries polling every **P** time units
- until the timeout **T** expires
- optionally with a context

# Ginkgo suite zoom in

```
var _ = Describe("Some foo cases", func() {  
    var x obj  
  
    BeforeEach(func() {  
        x = initialize()  
    })  
    Context("With some conditions", func() {  
        It("should behave like this", func() {  
            Expect(foo.Something()).To(Equal(somethingElse))  
        })  
    })  
    When("some other conditions apply", func() {  
        It("should behave like that", func() {  
            Expect(bar).Should(BeTrue())  
            Expect(baz).ToNot(BeNil())  
        })  
    })  
})
```



# Ginkgo suite zoom in

```
var _ = Describe("Some foo cases", func() {  
    BeforeEach(func() {  
        initialize()  
    })  
    Context("With some conditions", func() {  
        It("should behave like this", func() {  
            Expect(foo.Something()).To(Equal(somethingElse))  
        })  
    })  
    When("some other conditions apply", func() {  
        It("should behave like that", func() {  
            Expect(bar).Should(BeTrue())  
            Expect(baz).ToNot(BeNil())  
        })  
    })  
})
```

## Ginkgo suite zoom in

```
var _ = Describe "Some foo cases", func() {  
    BeforeEach(func() {  
        initialize()  
    })  
    Context("With some conditions", func() {  
        It("should behave like this", func() {  
            Expect(foo.Something()).To(Equal(somethingElse))  
        })  
    })  
    When "some other conditions apply", func() {  
        It("should behave like that", func() {  
            Expect(bar).Should(BeTrue())  
            Expect(baz).ToNot(BeNil())  
        })  
    })  
})
```

# Ginkgo suite zoom in

```
var _ = Describe("Some foo cases", func() {  
    BeforeEach(func() {  
        initialize()  
    })  
    Context("With some conditions", func() {  
        It("should behave like this", func() {  
            Expect(foo.Something()).To(Equal(somethingElse))  
        })  
    })  
    When("some other conditions apply", func() {  
        It("should behave like that", func() {  
            Expect(bar).Should(BeTrue())  
            Expect(baz).ToNot(BeNil())  
        })  
    })  
})
```

# Ginkgo suite zoom in

```
var _ = Describe("Some foo cases", func() {  
    BeforeEach(func() {  
        initialize()  
    })  
    Context("With some conditions", func() {  
        It("should behave like this", func() {  
            Expect(foo.Something()).To(Equal(somethingElse))  
        })  
    })  
    When("some other conditions apply", func() {  
        It("should behave like that", func() {  
            Expect(bar).Should(BeTrue())  
            Expect(baz).ToNot(BeNil())  
        })  
    })  
})
```

# Ginkgo suite zoom in

```
var _ = Describe("Some foo cases", func() {  
    BeforeEach(func() {  
        initialize()  
    })  
    Context("With some conditions", func() {  
        It("should behave like this", func() {  
            Expect(foo.Something()).To(Equal(somethingElse))  
        })  
    })  
    When("some other conditions apply", func() {  
        It("should behave like that", func() {  
            Expect(bar).Should(BeTrue())  
            Expect(baz).ToNot(BeNil())  
        })  
    })  
})
```

# How ginkgo runs

**ginkgo specs (= testcases) must be independent**

*repeating the operation multiple times produces the same result as applying it once, without any additional side effects*

ginkgo specs runs in random order and by default in parallel

“declare in container nodes, initialize in setup nodes”

# How ginkgo runs: walking the tree

Ginkgo runs in two steps: tree construction and run phase

## Tree Construction:

- Ginkgo visits all container nodes, invokes their closures and constructs the spec tree.
- Ginkgo captures the relevant setup and subject node closures by visiting the tree, but **does not run them**.

# How ginkgo runs: running the tree

Ginkgo runs in two steps: tree construction and run phase

Run phase:

- Ginkgo runs through each spec in the generated spec list sequentially.
- Ginkgo invokes the setup and subject nodes closures in the correct order and tracks any failed assertions, for each spec.
- Container node closures are never invoked.



## Common gotchas

- All Ginkgo nodes must only appear at the top-level or within a container node.
- A subject node cannot be top level

Note: you **CAN** nest arbitrarily container nodes though!

Note: you **CAN** have multiple top-level container nodes!

## Common gotchas /2

No assertion in container nodes! (ginkgo.Expect() ...)

Note: you **CAN** have any amount of assertions in a subject node!

# Common gotchas /3

## **Do not initialize variables in container nodes**

Subject nodes can mutate the values and pollute the state!

Perform initialization in setup nodes: these nodes are guaranteed to be called before every relevant subject node

Note: kinda OK for constants though - but should those be container variables?

# Logging: GinkgoWriter

GinkgoWriter is a globally available `io.Writer`.

Aggregates everything, only emits to stdout if the test fails.

`GinkgoWriter.TeeTo(writer)`: attach additional writers. Any data written to `GinkgoWriter` will immediately be sent to attached tee writers.

In verbose mode (`ginkgo -v`) writes to `GinkgoWriter` are immediately sent to stdout.

Logging: By() clause

By(“my message”)


Display the messages on failure

In verbose mode, displays the steps immediately

# Validating Webhooks

---

# Webhooks



# Validating Webhooks

- Bring your own validation logic
- Last resort, when the validation logic can't be implemented with kubebuilder annotations
- Better than having to look into logs!



# Validating Webhooks

```
kubebuilder create webhook --group workshop --version v1alpha1 --kind Configuration --programmatic-validation
```

```
INFO Writing kustomize manifests for you to edit...  
INFO Writing scaffold for you to edit...  
INFO internal/webhook/v1alpha1/configuration_webhook.go  
INFO internal/webhook/v1alpha1/configuration_webhook_test.go  
INFO internal/webhook/v1alpha1/webhook_suite_test.go  
INFO Update dependencies
```

# Validating Webhooks

```
// ValidateCreate implements webhook.CustomValidator so a webhook will be registered for the type
Configuration.
func (v *ConfigurationCustomValidator) ValidateCreate(_ context.Context, obj runtime.Object)
(admission.Warnings, error) {
    configuration, ok := obj.(*workshopv1alpha1.Configuration)
    if !ok {
        return nil, fmt.Errorf("expected a Configuration object but got %T", obj)
    }
    configurationlog.Info("Validation for Configuration upon creation", "name",
configuration.GetName())

    // TODO(user): fill in your validation logic upon object creation.

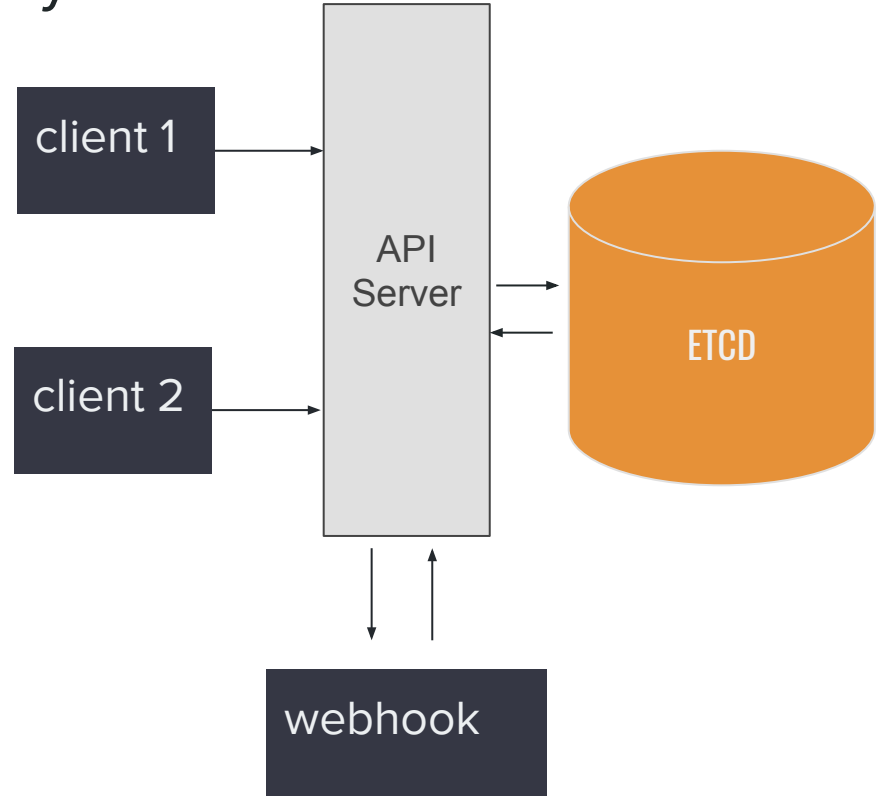
    return nil, nil
}
```

## Use webhooks when

- You need to validate multiple objects compatibility
- The validation is complicated and goes beyond the syntax
- You need to validate against the internal state

# Beware of eventual consistency!

- Multiple clients might perform multiple calls against the api server
- Having a webhook does not guarantee consistency
- **Always** validate in your reconcile loop (possibly using the same logic)



# Wrapping Up

---

Distributed systems are  
complicated

Failure is normal

Corner cases are hard to  
handle



Edge driven is tempting, level  
driven is more robust

Controller Runtime gives us a  
safe framework

# Resources

- The kubebuilder book: [book.kubebuilder.io](https://book.kubebuilder.io)
- Kubernetes custom resources:  
[kubernetes.io/docs/tasks/extend-kubernetes/custom-resources/custom-resource-definitions/](https://kubernetes.io/docs/tasks/extend-kubernetes/custom-resources/custom-resource-definitions/)

Thank you!

---